Research Statement

Overview

My research explores fundamental, everyday decisions about individuals' daily lives, especially about work and family. I am interested in family structures: why and when individuals decide to partner or have children; how they organize market work, housework, and caregiving tasks; and how these decisions are affected by societal forces including gender norms and public policy. While these decisions are ubiquitous, choices about work, partnering, and children have major impacts on individual well-being. Moreover, research increasingly shows that these household-level decisions are central to explaining most of the remaining gender inequality in labor outcomes in high-income countries (Kleven, Landais, and Søgaard, 2019; Angelov, Johansson, and Lindahl, 2016; Adda, Dustmann, and Stevens, 2017; Olivetti and Petrongolo, 2017).

As an applied microeconomist, I have skills in modern causal inference, particularly with respect to the latest work on difference-in-difference estimators and shift-share instruments, both of which I have used in existing projects. However, my work is not solely focused on identifying causal effects. My research also employs descriptive empirics and theoretical models, each of which has been essential for key insights in my work thus far.

Within the overarching motivation of studying family structures, my research agenda can be grouped roughly into three interlinking subtopics: partnering decisions, gender norms, and reproductive choices and outcomes. For each of these topics, I will describe the projects I have carried out during my PhD, as well as planned and early-stage projects going forward.

Partnering Decisions

Partnering decisions are at the core of understanding family structures. People choose whether to cohabit or marry (or to separate or remain single) based on a number of factors, with some of the most important factors being whether a relationship provides love and companionship, that it can reduce living costs and the costs of raising children, and that members of a household can insure each other against negative shocks. I am particularly interested in understanding how exposure to or protection from financial risk affects partnering decisions, especially whether it can affect the quality of relationships that arise.

Theory and empirics are both essential in studying partnering decisions, and my job market paper, *Public health insurance and marital outcomes: Evidence from the Medicaid expansion* exemplifies my experience in combining the two. Co-authored with Tom Potoms, a fellow PhD at the time we began, this project explores how public insurance affects marital outcomes through the lens of Medicaid expansion in the United States. We develop a search model of the marriage market with transferable utility in households, which predicts that expansion of public insurance will lead to an increase in match quality without necessarily leading to an increase in divorce. Empirically, we show that increases in the likelihood of eligibility for Medicaid due to policy expansions leads to lower marriage rates for individuals in their 20s and 30s, and lower divorce rates for new marriages, both of which are consistent with increased marital match quality. A novel finding of our results is that marriage rates decrease even for

individuals with high-educational attainment, suggesting means-tested programs may have broader behavioral responses than has sometimes been considered. Together, these results show that public insurance can meaningfully affect marriage outcomes, even for groups we might not expect to respond to the introduction of a means-tested program.

In a similar spirit, with this same co-author and an additional PhD student, we are also embarking on a new project studying how student loans in the U.S. affect marriage, cohabitation, and marital stability. Student loans comprise the second largest category of debt in the U.S. after mortgages, yet in contrast to mortgages, they are typically incurred prior to marriage. In a way, this project represents the flipside of public insurance: student loans represent an increase in financial burden, even as degrees obtained as a result of student loans may yield increased earnings. Existing work has found evidence of short-run decreases in marriage rates as a result of increased student loans. Our aim is to develop and estimate a structural model of cohabitation, marriage, and labor supply to study the longer-run impacts of student loans on family structures.

Gender Norms

Gender norms broadly defined have been shown to be negatively related to female labor force participation (Fortin, 2005; Fernandez, 2007; Fortin, 2015; Goussé, Jacquemet, and Robin, 2017), but precisely what norms matter and how they operate is still less well-understood. The motivation of the first paper of my PhD, *Revisiting the breadwinner norm: The effect of the potential relative wage on married women's labor supply* (under a second round of revision at JOLE), was to hone in on one particular norm, known as the "male breadwinner norm". This project highlights my skills in using descriptive and non-causal empirical analysis to generate insights. Specifically, I aimed to understand whether an aversion to female breadwinning affects women's labor supply decisions within marriage, as prior results have suggested, or whether there was another possible explanation for the findings that women appeared to be less likely to work after out-earning their husbands, which were based on measures of observed relative earnings.

In examining characteristics associated with female breadwinning, I show that the identifying variation offered by observed relative earnings is not informative for understanding the response to wives' earning more in itself because realized female breadwinning largely coincides with low levels of work for husbands. In studying the effects of observed earnings, we cannot distinguish whether our estimates are primarily owe to men's reductions in work or to female breadwinning itself. When I instead use a measure of potential relative earnings based on average earnings for demographic groups, which is clearly linked to the likelihood of female breadwinning but not men's reductions in work, I find no effect of an increased likelihood of female breadwinning on wives' labor force participation. To the extent that there is an aversion to female breadwinning, this result implies that it does not cause reductions in married women's labor force participation. Although the results of this first project imply no effects of an aversion to female breadwinning *within marriage*, it does not rule out that this norm may affect the earlier stage of partnering decisions.

The implication that this norm might still affect partnering decisions has led me to consider more broadly what norms might be important for partnering decisions and how. A different way of conceptualizing the male breadwinner norm is that it is important for men to be working full-time, rather than to be earning more than their wives. Indeed, whether or not we call it a norm, this tendency is a striking empirical regularity: it is well-established as an empirical fact that men have low wage elasticities in high-income countries, and when observed to be not working or working part-time, are almost always involuntarily under- or unemployed.

The next planned project in my agenda would explore how this empirical fact or norm for affects marriage markets, and in turn, key decisions that contribute to gender inequality over the life course, which would make use of both my experience studying partnering decisions and gender norms. A result of the empirical fact that men always work (or always intend to be working) is that heterosexual women face a market of partners who they can typically expect to share at most half of domestic and childcare duties (for example, if both partners work full-time), whereas heterosexual men in contrast face a market where they can expect potential partners to do at minimum half of domestic and childcare duties. This simple asymmetry in expectations regarding expectations about partners could theoretically have impacts on educational and occupational choices, as well as selection into partnerships and parenthood, and ultimately, gender inequality. One challenge with implementing this project is that there is limited empirical variation: although there is variation in the average hours of housework men carry out, there are really no countries that have more than a negligible share of men that work part-time or leave the workforce to be primary caregivers, as large shares of women nearly everywhere do. Even despite this limitation, this framework could serve to help explain why policies that incentivize women's labor force attachment after having children (such as maternity leave and childcare subsidies) have little effect on the within-household child-penalty: choices about the (future) division of labor may be to a large extent already determined at the time of couple formation.

Reproductive Choices and Outcomes

Children are an undercurrent of questions of about partnering decisions, gender norms, and gender inequality: for partnering decisions, they are an important reason why many couples form, and people sort along preferences for children. Many if not most aspects of gender norms are with respect to gendered roles in caring for and raising children. As previously highlighted, the child penalty explains most remaining gender inequality in high income countries. Understanding the decisions to have children and how to raise them are thus natural extensions of these topics.

The third chapter of my PhD thesis, *Subsidizing contraception: Effects on takeup of highly effective contraceptives and unintended births*, is motivated by the fact about a third of births in the U.S. are still unintended, meaning mothers would have preferred that the pregnancies would have occurred at least two years later or not at all (Buckles, Guldi, and Schmidt, 2019). At the same time, over the past 10 years, the rate of unintended births has substantially decreased. I explore whether subsidization of prescription contraceptives has contributed to

this decline. I use state-of-the-art techniques for estimation of difference-in-differences to evaluate how means-tested public health insurance programs under the umbrella of Medicaid affected take-up of highly effective contraceptives and in turn, unintended fertility. While prior estimates of some of the earliest programs found the subsidization of contraceptives resulted in lowered fertility (DeLeire, Lopoo, and Simon, 2011; Kearney and Levine, 2009), I find no evidence for reductions in fertility for later programs, despite that prescriptions for highly effective contraceptives still increase.

An advantage of using the most recent techniques in differences-in-differences estimation is that they allow for exploring multiple dimensions of treatment effect heterogeneity, and in this case, doing so illustrates its importance: while these programs look comparable across states in terms of eligibility thresholds, some states exhibit no reductions in fertility while others exhibit substantial reductions. This heterogeneity implies that providing subsidies by themselves may not be enough to affect contraceptive choices, and in turn fertility outcomes, and that these programs may differ by other meaningful characteristics. This insight is important both for considering how policy relates to unintended births, as well as interpreting estimates of the effects of Medicaid programs in other dimensions, where there may well be similar types of underlying policy heterogeneity.

A final project in a preliminary stage would explore the long-run impacts of breastfeeding owing to changes in hospital policies in Sweden between the 1970s and 1990s, with Dan-Olof Rooth at Stockholm University. Breastfeeding is one issue where the consensus among the medical community and policymakers is overwhelmingly in favor, although the evidence for its beneficial health effects is in fact modest. Moreover, the existing literature evaluates effects up until about age 10, so this project could provide a valuable contribution in assessing whether there are long-run benefits.

The hospital policy changes of interest span 90 institutions over a period of 35 years, and as such documentation is not easily found. I have implemented a novel method to aid in identifying when relevant changes occurred, by recruiting participants on Facebook to a short, anonymous, non-representative survey of individuals' experiences when their children were born. This method is extremely cheap and fast, resulting in approximately 7,000 responses over a two-week period that allow me to pinpoint roughly when changes occurred at a given hospital, which will aid in narrowing down archival searches describing changes in greater detail. This project has remained at this stage for the past year, since the pandemic made it infeasible to travel around Sweden visiting regional archives, but this should likely not be a limitation in the upcoming year.

1 References

Adda, J., C. Dustmann, and K. Stevens (2017). The career costs of children. *Journal of Political Economy 125*(2), 293–337.

Angelov, N., P. Johansson, and E. Lindahl (2016). Parenthood and the gender gap in pay. *Journal of Labor Economics 34*(3), 545–579.

Buckles, K., M. E. Guldi, and L. Schmidt (2019). Fertility trends in the united states, 1980-2017: The role of unintended births. Technical report, National Bureau of Economic Research.

DeLeire, T., L. M. Lopoo, and K. I. Simon (2011). Medicaid expansions and fertility in the united states. *Demography 48*(2), 725–747.

Fernandez, R. (2007). Women, work, and culture. *Journal of the European Economic Association 5*(2-3), 305–332.

Fortin, N. M. (2005). Gender role attitudes and the labour-market outcomes of women across OECD countries. *Oxford Review of Economic Policy 21*(3), 416–438.

Fortin, N. M. (2015). Gender role attitudes and women's labor market participation: Opting-out, aids, and the persistent appeal of housewifery. *Annals of Economics and Statistics* (117-118), 379–401.

Goussé, M., N. Jacquemet, and J.-M. Robin (2017). Marriage, labor supply, and home production. *Econometrica 85*(6), 1873–1919.

Kearney, M. S. and P. B. Levine (2009). Subsidized contraception, fertility, and sexual behavior. *The review of Economics and Statistics 91*(1), 137–151.

Kleven, H., C. Landais, and J. E. Søgaard (2019). Children and gender inequality: Evidence from Denmark. *American Economic Journal: Applied Economics 11*(4), 181–209.

Olivetti, C. and B. Petrongolo (2017). The economic consequences of family policies: lessons from a century of legislation in high-income countries. *Journal of Economic Perspectives 31*(1), 205–30.